2. ec2-34-200-229-125.compute-1.ama

```
Authenticating with public key "Imported-Openssh-Key"

                • MobaXterm Personal Edition v26.3 •
              (SSH client, X server and network tools)

 ➤ SSH session to hadoop@ec2-34-200-229-125.compute-1.amazonaws.com
   • Direct SSH      :  ✔
   • SSH compression :  ✔
   • SSH-browser     :  ✔
   • X11-forwarding  :  ✘  (disabled or not supported by server)

 ➤ For more info, ctrl+click on help or visit our website.


Last login: Sat May  9 06:46:48 2026

       __|  __|_  )
       _|  (     /   Amazon Linux AMI
      ___|\___|___|

https://aws.amazon.com/amazon-linux-ami/2018.03-release-notes/
49 package(s) needed for security, out of 74 available
Run "sudo yum update" to apply all updates.

EEEEEEEEEEEEEEEEEEEE MMMMMMMM           MMMMMMMM RRRRRRRRRRRRRRR
E::::::::::::::::::E M:::::::M         M:::::::M R::::::::::::::R
EE:::::EEEEEEEEE:::E M::::::::M       M::::::::M R:::::RRRRRR:::::R
  E::::E       EEEEE M:::::::::M     M:::::::::M RR::::R      R::::R
  E::::E             M::::::M:::M   M:::M::::::M   R:::R      R::::R
  E:::::EEEEEEEEEE   M:::::M M:::M M:::M M:::::M   R:::RRRRRR:::::R
  E::::::::::::::E   M:::::M  M:::M:::M  M:::::M   R::::::::::RR
  E:::::EEEEEEEEEE   M:::::M   M:::::M   M:::::M   R:::RRRRRR::::R
  E::::E             M:::::M    M:::M    M:::::M   R:::R      R::::R
  E::::E       EEEEE M:::::M     MMM     M:::::M   R:::R      R::::R
EE:::::EEEEEEEE::::E M:::::M             M:::::M   R:::R      R::::R
E::::::::::::::::::E M:::::M             M:::::M RR::::R      R::::R
EEEEEEEEEEEEEEEEEEEE MMMMMMM             MMMMMMM RRRRRRR      RRRRRR

[hadoop@ip-172-31-5-58 ~]$ hive
```

monitoring

```
[hadoop@ip-172-31-5-58 ~]$ hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: false
hive> CREATE TABLE crime_data (
    >     crime_id INT,
    >     case_number STRING,
    >     crime_date STRING,
    >     block STRING,
    >     primary_type STRING,
    >     description STRING,
    >     location_desc STRING,
    >     arrest STRING,
    >     domestic STRING,
    >     beat INT,
    >     district INT,
    >     ward INT,
    >     community_area INT,
    >     fbi_code INT,
    >     latitude DOUBLE,
    >     longitude DOUBLE
    > )
    > ROW FORMAT DELIMITED
    > FIELDS TERMINATED BY ','
    > STORED AS TEXTFILE;
OK
Time taken: 2.758 seconds
hive> show tables;
OK
crime_data
Time taken: 0.168 seconds, Fetched: 1 row(s)
```

```
hive> CREATE TABLE crimes_partitioned (
    >     crime_id INT,
    >     case_number STRING,
    >     crime_date STRING,
    >     block STRING,
    >     primary_type STRING,
    >     description STRING,
    >     location_desc STRING,
    >     arrest STRING,
    >     domestic STRING,
    >     beat INT,
    >     district INT,
    >     ward INT,
    >     community_area INT,
    >     fbi_code INT,
    >     latitude DOUBLE,
    >     longitude DOUBLE
    > )
    > PARTITIONED BY (crime_year INT)
    > CLUSTERED BY (district)
    > INTO 4 BUCKETS
    > ROW FORMAT DELIMITED
    > FIELDS TERMINATED BY ','
    > STORED AS TEXTFILE;
OK
Time taken: 0.11 seconds
hive> SELECT COUNT(*) AS total_crimes
    > FROM crime_partitioned;
FAILED: SemanticException [Error 10001]: Line 2:5 Table not found 'crime_partitioned'
hive> SELECT COUNT(*) AS total_crimes
    > FROM crimes_partitioned;
OK
0
Time taken: 2.87 seconds, Fetched: 1 row(s)
hive> SELECT primary_type,
    >        COUNT(*) AS total_cases
    > FROM crimes_partitioned
```

```
    > FROM crimes_partitioned;
OK
0
Time taken: 2.87 seconds, Fetched: 1 row(s)
hive> SELECT primary_type,
    >        COUNT(*) AS total_cases
    > FROM crimes_partitioned
    > GROUP BY primary_type
    > ORDER BY total_cases DESC
    > LIMIT 5;
Query ID = hadoop_20260509071458_07ba15ce-465b-4d98-bcef-b6b39e37caae
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1778308816614_0002)

--------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
--------------------------------------------------------------------------------
Map 1            container     SUCCEEDED      0          0        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      1          1        0        0       0       0
Reducer 3 ...... container     SUCCEEDED      1          1        0        0       0       0
--------------------------------------------------------------------------------
VERTICES: 02/03  [==========================>>] 100%  ELAPSED TIME: 5.38 s
--------------------------------------------------------------------------------
OK
Time taken: 15.521 seconds
hive> SELECT crime_year,
    >        COUNT(*) AS total_crimes
    > FROM crimes_partitioned
    > GROUP BY crime_year
    > ORDER BY crime_year ASC;
Query ID = hadoop_20260509071532_172f7071-55d1-4ba1-9164-cbd210b7e5f4
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1778308816614_0002)
```

```
Query ID = hadoop_20260509071532_172f7071-55d1-4ba1-9164-cbd210b78514
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1778308816614_0002)

--------------------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
--------------------------------------------------------------------------------------------
Map 1            container     SUCCEEDED      0          0        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      1          1        0        0       0       0
Reducer 3 ...... container     SUCCEEDED      1          1        0        0       0       0
--------------------------------------------------------------------------------------------
VERTICES: 02/03  [==========================>>] 100%  ELAPSED TIME: 4.15 s
--------------------------------------------------------------------------------------------
OK
Time taken: 5.213 seconds
hive> SELECT primary_type,
    >        COUNT(*) AS total_cases,
    >        SUM(CASE WHEN arrest = 'true' THEN 1 ELSE 0 END) AS total_arrests,
    >        (SUM(CASE WHEN arrest = 'true' THEN 1 ELSE 0 END) * 100.0 / COUNT(*)) AS arrest_rate
    > FROM crimes_partitioned
    > GROUP BY primary_type
    > HAVING SUM(CASE WHEN arrest = 'true' THEN 1 ELSE 0 END) > 500
    > ORDER BY arrest_rate DESC;
Query ID = hadoop_20260509071604_d3671cbd-120d-44c0-926d-12be2f96a075
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1778308816614_0002)

--------------------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
--------------------------------------------------------------------------------------------
Map 1            container     SUCCEEDED      0          0        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      1          1        0        0       0       0
Reducer 3 ...... container     SUCCEEDED      1          1        0        0       0       0
--------------------------------------------------------------------------------------------
VERTICES: 02/03  [==========================>>] 100%  ELAPSED TIME: 3.73 s
--------------------------------------------------------------------------------------------
OK
Time taken: 4.871 seconds
```

```
OK
Time taken: 4.871 seconds
hive> SELECT *
    > FROM crimes_partitioned
    > WHERE crime_year = 2020;
OK
Time taken: 0.664 seconds
hive> SELECT district,
    >        COUNT(*) AS domestic_crime_count
    > FROM crimes_partitioned
    > WHERE domestic = 'true'
    > GROUP BY district
    > ORDER BY domestic_crime_count DESC
    > LIMIT 3;
Query ID = hadoop_20260509071742_d3887cb4-2a22-4f05-a815-26948bd5a09b
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1778308816614_0002)

--------------------------------------------------------------------------------
        VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
--------------------------------------------------------------------------------
Map 1            container     SUCCEEDED      0          0        0        0       0       0
Reducer 2 ...... container     SUCCEEDED      1          1        0        0       0       0
Reducer 3 ...... container     SUCCEEDED      1          1        0        0       0       0
--------------------------------------------------------------------------------
VERTICES: 02/03  [==========================>>] 100%  ELAPSED TIME: 3.57 s
--------------------------------------------------------------------------------
OK
```

with other AWS services.

- **Choose an existing service role** — Select a default service role or a custom role with IAM policies attached so that your cluster can interact with other AWS services.
- **Create a service role** — Let Amazon EMR create a new service role so that you can grant and restrict access to resources in other AWS services.

Service role

EMR_DefaultRole

EC2 instance profile for Amazon EMR

The instance profile assigns a role to every EC2 instance in a cluster. The instance profile must specify a role that can access the resources for your steps and bootstrap actions.

- **Choose an existing instance profile** — Select a default role or a custom instance profile with IAM policies attached so that your cluster can interact with your resources in Amazon S3.
- **Create an instance profile** — Let Amazon EMR create a new instance profile so that you can specify a custom set of resources for it to access in Amazon S3.

Instance profile

EMR_EC2_DefaultRole

Custom automatic scaling role - *optional*

When a custom automatic scaling rule triggers, Amazon EMR assumes this role to add and terminate EC2 instances.

Learn more

Custom automatic scaling role

Summary Info

Name and applications - *required*

Name
newcluster

Amazon EMR release
emr-5.20.1

Application bundle
Custom (Ganglia 3.7.2, Hadoop 2.8.5, Hive 2.3.4, Hue 4.3.0, JupyterHub 0.9.4, Spark 2.4....)

Cluster configuration - *required*

Uniform instance groups
Primary (m4.large), Core (m4.large), Task (m5.xlarge)

Cluster scaling and provisioning - *required*

Provisioning configuration
Core size: 1 instance
Task size: 1 instance

Cancel Create cluster

Amazon EMR > EMR on EC2: Clusters > newcluster

Your cluster "newcluster" has been successfully created.

Notifications 0 0 2 0 0

newcluster

Updated 8 minutes ago | Terminate | Clone in AWS CLI | Clone

This EMR release reaches End of Support on Aug-01-2026 and will no longer be eligible for technical support. AWS strongly recommends that you run your workloads on the latest Amazon EMR release to receive security-critical updates and fixes. You can also use new Spark upgrade agent to modernize your existing applications on version 5.40 or higher to latest EMR version. To learn more, see EMR Standard Support policy and Spark Upgrades.

Summary

Cluster info

Cluster ID
j-3SFA3VTIN6VZY

Cluster ARN
arn:aws:elasticmapreduce:us-east-1:339712764081:cluster/j-3SFA3VTIN6VZY

Cluster configuration
Instance groups

Capacity
1 Primary | 1 Core | 1 Task

Applications

Amazon EMR version
emr-5.20.1

Installed applications
Ganglia 3.7.2, Hadoop 2.8.5, Hive 2.3.4, Hue 4.3.0, JupyterHub 0.9.4, Spark 2.4.0, Tez 0.9.1, Zeppelin 0.8.0

Cluster management

Log destination in Amazon S3
aws-logs-339712764081-us-east-1/elasticmapreduce

Primary node public DNS
ec2-34-200-229-125.compute-1.amazonaws.com

Connect to the Primary node using SSH

Connect to the Primary node using SSM

Status and time

Status
Waiting

Creation time
May 09, 2026, 11:58 (UTC+05:30)

Elapsed time
24 minutes

Properties | Bootstrap actions | Instances (Hardware) | Steps | Applications | Configurations | Monitoring | Events | Tags (0)

Query

Search saved documents...

Jobs yash

ault

(0)

s found

File Browser

Search for file name

Actions

Delete forever

Upload

New

Home / user / hadoop

	Name	Size	User	Group	Permissions	Date
	↑		hdfs	hadoop	drwxr-xr-x	May 08, 2026 11:53 PM
	.		hadoop	hadoop	drwxrwxrwx	May 08, 2026 11:56 PM
	Crimes_-_2022_20260507 - Crimes_-_2022_20260507.csv	54.2 MB	yash	hadoop	-rw-r--r--	May 08, 2026 11:56 PM

Show 45 of 1 items

Page 1 of 1